

Surveillance Poisons Communication in LLM Global Games

Khaled Eltokhy
Department of Economics
The Graduate Center, CUNY

April 2026

Abstract

Authoritarian regimes can suppress coordinated political action either by directly deterring would-be dissenters or by poisoning the channels through which they communicate. The two mechanisms look similar in field data, where institutional and selection differences confound everything. I isolate the second using LLM agents in the Morris–Shin (2003) regime-change game: in clean prompt-isolation reruns, appending a monitoring warning to the *communication* prompt—leaving the *decision* prompt unchanged—lowers joining in 5/5 models (mean $\Delta = -8.4$ pp). Placebo (research-only) and anonymous monitoring warnings do not replicate the effect. A discriminant test (Llama 3.3 70B) decomposes the channel: surveillance moves a *non-coordinative* version of the same decision (a private bet on regime fall, payoff independent of others) by 5.6 pp—an informational degradation that affects any decision drawing on peer messages—while the coordination-task delta in the same model is 9.8 pp; the residual 4.2 pp is a coordination-specific component, detectable only when peer messages serve as coordination cues. A logistic classifier distinguishes baseline from surveillance messages at 91.9% accuracy in the primary model (held-out $N = 15000$; replicating at 97–99% accuracy in Llama 3.3 70B and Qwen3 30B): the shift is from direct (“regime,” “security forces,” “streets”) to coded (“walls cracking,” “ground shifting,” “heads down”) language, with regime-weakness references actually rising under monitoring in the primary model. Across 1,800 period-level observations the agents also display threshold-like policies aligned with the theoretical attack-mass benchmark (mean $r = +0.80$, $p < 0.001$ for every model); this correlation survives every treatment, indicating level shifts rather than disrupted strategic logic. The same channel that helps citizens coordinate is the channel an authoritarian regime degrades.

JEL: C72, C92, D82, D83, P16

Keywords: global games, regime change, surveillance, communication, LLM agents, preference falsification

1 Introduction

The information channel is a trap. In the canonical global game of regime change (Morris and Shin, 2003), private

signals alone resolve equilibrium multiplicity (Carlsson and van Damme, 1993; Frankel et al., 2003)—but real coordination crises, from bank runs (Diamond and Dybvig, 1983) to currency attacks (Obstfeld, 1996) to political upheaval (Angeletos et al., 2007), also feature communication among agents. Any channel that transmits information about others’ willingness to act is also the channel through which regimes suppress it. This paper augments the Morris–Shin framework with an explicit pre-play communication network and documents how surveillance exploits that channel.

The idea behind a global game is simple. A regime has some underlying strength, but no citizen knows it exactly. Each citizen receives a private, noisy signal—some see evidence of weakness, others see evidence of stability—and must decide whether to join an uprising or stay home. Joining is risky: it pays off only if enough others join too. Staying home is safe but forgoes the chance of change. The strategic tension is that each citizen’s best action depends on what they believe *others* will do, which in turn depends on what others believe, and so on. Carlsson and van Damme (1993) and Morris and Shin (2003) showed that when citizens receive sufficiently precise private signals, this infinite regress collapses to a unique equilibrium: there is a single threshold above which citizens stay and below which they join. This result transformed the study of coordination failures—currency crises, bank runs, political revolutions—by providing a tractable framework with sharp comparative statics.

Testing these predictions empirically is difficult. Laboratory experiments confirm threshold behavior with stylized numeric signals, but cannot manipulate the rich, multi-source information that characterizes real crises. Field data from actual coordination failures confounds strategic behavior with institutional differences and unobserved information structures. The gap between the theory’s tractability and the difficulty of testing it empirically motivates this paper’s approach.

The empirical strategy treats LLM agents as a measurement instrument, useful insofar as they reproduce known theoretical benchmarks before treatment effects are identified. Seven models spanning six architecture families are embedded in the Morris and Shin (2003) game; each receives a private signal $x_i = \theta + \varepsilon_i$ rendered as a

natural-language intelligence briefing whose content shifts monotonically with signal strength across eight evidence domains. No explicit payoff table is provided—stakes are embedded in the narrative, forcing agents to form beliefs from language. Section 4 establishes the behavioral benchmark; the surveillance treatment is introduced only after threshold alignment is verified (Section 5).

Contribution roadmap. The paper’s core contribution is the surveillance result. I show that the communication channel that enables coordination is also exploitable by an authoritarian planner. In clean prompt-isolation reruns, surveillance lowers joining in 5/5 models (mean $\Delta = -8.4$ pp; primary-model $\Delta = -9.1$ pp), while placebo and anonymous monitoring warnings leave primary-model behavior unchanged. A degraded-messages control and a no-message control both reduce joining, consistent with the view that decision-stage agents use peer messages as part of their information environment. A discriminant test in Llama 3.3 70B decomposes the channel: surveillance moves a *non*-coordinative version of the same decision (a private bet on regime fall, payoff independent of others) by 5.6 pp—an informational degradation any peer-message-using decision would inherit—while the coordination-task delta in the same model is 9.8 pp. The residual 4.2 pp is a coordination-specific component, detectable only when peer messages serve as coordination cues. Surveillance suppresses coordination above and beyond its informational effect. Elicited beliefs shift under surveillance, but actions shift further; the residual wedge is reported as a measured quantity, not a mechanism claim.

The surveillance claim rests on a behavioral foundation. Across 1,800 period-level observations and seven models, LLM agents display robust threshold-like behavior aligned with the global-game benchmark (mean $r = +0.80$, $p < 0.001$ for every model); a narrative stakes ladder confirms that LLMs respond to strategic framing, not just text sentiment (Appendix E.10).

Related literature. The paper connects two strands. The global game literature—from Carlsson and van Damme (1993) through Morris and Shin (2003) and the dynamic extensions of Angeletos et al. (2007) and Chassang (2010)—provides the theoretical framework; laboratory experiments (Heinemann et al., 2004, 2009; Szkup and Trevino, 2020) confirmed threshold behavior with numeric signals. The political economy of authoritarian control provides the substantive motivation. Kuran (1991) models preference falsification under social pressure; Edmond (2013) formalizes propaganda’s effect on coordination beliefs; Shadmehr and Bernhardt (2011) characterizes the collective action problem under repression. Empirically, King et al. (2013) document that Chinese censorship targets posts about collective action specifically—suggesting authorities perceive the coordination channel as the binding threat. Penney (2016) and Stoycheff (2016) show surveillance announcements suppress online expres-

sion; Enikolopov et al. (2020) that social media facilitates protest. These literatures establish a robust empirical fact—surveillance and censorship reduce dissident behavior—but field data cannot separate the two channels through which they could operate: *direct deterrence* (citizens fearing punishment for participation) and *communication poisoning* (peers self-censoring such that downstream agents lose coordination cues). The contribution here is not to re-establish that surveillance suppresses dissent—field work has done that—but to isolate the second channel using a design that holds individual deterrence fixed and varies only the messaging environment.¹

On the methodological side, Horton (2023) proposed LLMs as “homo silicus.” Subsequent work found LLMs struggle in coordination games (Akata et al., 2025; Sun et al., 2025), with model scale not predicting strategic performance (Petrov et al., 2025). Huang et al. (2024) and Carlini et al. (2025) document that alignment shifts risk preferences, motivating narrative rather than tabular payoffs. Critical reviews by Larooij and Törnberg (2025a,b) identify validation challenges that I address through cross-generator robustness (Appendix E.7), scramble/flip falsification, and out-of-sample replication.

The paper proceeds as follows. Section 2 presents the theoretical framework. Section 3 describes the experimental design. Section 4 establishes a threshold-like behavioral benchmark and shows that pre-play communication creates an informative channel. Section 5 tests whether an authoritarian regime can exploit that channel through surveillance. Section 6 concludes.

2 The Global Game of Regime Change

A continuum of citizens indexed by $i \in [0, 1]$ simultaneously choose whether to join an uprising ($a_i = 1$) or stay home ($a_i = 0$). The regime has strength $\theta \in \mathbb{R}$, drawn from a diffuse (improper uniform) prior. States $\theta \leq 0$ represent regimes so weak they fall without opposition; states $\theta \geq 1$ represent regimes that survive even unanimous attack. The regime falls if the mass of citizens who join exceeds θ (Eq. 1):

$$\text{Regime falls} \iff A \equiv \int_0^1 a_i di > \theta. \quad (1)$$

Payoffs depend on the citizen’s action and the outcome (Eq. 2):

$$u_i(a_i, A, \theta) = \begin{cases} B & \text{if } a_i = 1 \text{ and } A > \theta \\ -C & \text{if } a_i = 1 \text{ and } A \leq \theta \\ 0 & \text{if } a_i = 0 \end{cases} \quad (2)$$

¹Concurrent work by Ruano and Rajan (2026) places LLM depositors in a bank-run coordination game with network contagion, finding sharp phase transitions consistent with Diamond–Dybvig.

where $B > 0$ is the payoff to joining a successful uprising and $C > 0$ is the cost of joining a failed attempt. Non-participants receive zero regardless of the outcome.

Each citizen observes a private signal $x_i = \theta + \varepsilon_i$, where $\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$ independently across citizens.

Proposition 1 (Morris and Shin, 2003). *In the limit of diffuse priors, there exists a unique Bayesian Nash equilibrium in threshold strategies. An agent joins if and only if $x_i < x^*$ (Eq. 3), where*

$$x^* = \theta^* + \sigma \Phi^{-1}(\theta^*) \quad (3)$$

and $\theta^* = B/(B + C)$.

The *attack mass* $A(\theta)$ —the theoretical fraction of the population that joins at regime strength θ —is given by:

$$A(\theta) = \Phi\left(\frac{x^* - \theta}{\sigma}\right). \quad (4)$$

This is a decreasing function of θ : weaker regimes face larger uprisings.

I define $J(\theta)$ as the *empirical join fraction*—the finite-sample, observed counterpart to the theoretical attack mass $A(\theta)$. My experiment uses $n = 25$ agents with a normal prior $\theta \sim \mathcal{N}(\bar{z}, 1)$, and agents do not observe (B, C, σ) directly. I use the continuum $A(\theta)$ as an external behavioral benchmark throughout, not as the exact equilibrium of the implemented finite- n game.

The briefing generator—the system that converts z-scores into natural-language intelligence briefings—has three tunable control parameters (clarity, directional precision, dissent framing). Full generator details are in Appendix E.11.

3 Experimental Design

The experiment has two phases. First, I establish that LLM agents display threshold-like policies in the global game when private signals are conveyed in natural language: a pure treatment (private signals only), a communication treatment (pre-play messaging), and falsification tests. Second, I test whether an authoritarian regime can exploit the communication channel through surveillance. All LLM interactions use the same prompt structure across models.

In the theory (Section 2), θ is an unbounded fundamental. In the benchmark experiments, θ is drawn from a normal distribution and agents are not shown payoff parameters (B, C) . I compute the benchmark $A(\theta)$ under the canonical normalization $B = C = 1$ ($\theta^* = 0.50$, $\sigma = 0.3$).

Throughout, $\theta^* = B/(B + C)$ is the theoretical cutoff, $x^* = \theta^* + \sigma \Phi^{-1}(\theta^*)$ the signal cutoff, and $\hat{\theta}^*$ the estimated cutoff from the fitted logistic $P(\text{join} \mid \theta) = 1/(1 + e^{b_0 + b_1 \theta})$, where $b_1 > 0$ corresponds to a join curve that is *decreasing* in θ (the expected direction). The cutoff is $\hat{\theta}^* = -b_0/b_1$. I write $\beta \equiv b_1$ for the logistic steepness parameter in tables.

For each country–period, nature draws $\theta \sim \mathcal{N}(\bar{z}, 1)$, where $\bar{z}$ is a public prior mean drawn randomly for each country. Each agent i receives a private signal $x_i = \theta + \varepsilon_i$. The platform normalizes that signal to $z_i = (x_i - \bar{z})/\sigma$ and translates z_i into a multi-paragraph intelligence briefing via a deterministic generator that maps signal strength to narrative content about regime stability, economic conditions, public sentiment, and coordination prospects. Agents observe the briefing, not the normalization inputs $(\bar{z}, \sigma)$ themselves. Figure 1 summarizes the signal-to-text-to-decision pipeline.

Agents observe only their private briefing and never the payoff parameters (B, C) , the noise level σ , or the prior distribution. The diffuse-prior equilibrium (Proposition 1) serves as an external benchmark for behavioral shape and comparative statics, not as a claim about the exact equilibrium of the implemented game.

The behavioral benchmark has four core treatments: a *pure global game* (independent decisions), *communication* (pre-play messaging on a Watts–Strogatz network, $k = 4$, $p = 0.3$), *scramble* (cross-period briefing permutation breaking the θ -to-signal link), and *flip* (negated z-scores). A *degraded-messages* control replaces LLM messages with generic text, and a no-message control removes peer messages entirely; together they benchmark how much the decision stage depends on message content. Additional robustness tests are in Appendix E.

I test seven models spanning six families (Table 1). All experiments use $n = 25$ agents per period-level task row and $\sigma = 0.3$. Benchmark experiments use $B = C = 1$ ($\theta^* = 0.50$). All LLM calls use temperature = 0.7 with a single sample per decision, routed through OpenRouter and cached by request hash (Appendix F). Per-model calibration adjusts only a single scalar (`cutoff_center`, the logistic offset that maps z-scores to directional sentiment); all other generator parameters are fixed ex ante. Placebo perturbations of ± 0.3 to this scalar leave the threshold correlation virtually unchanged (Appendix E, Figure A5): calibration shifts the *level* of the join curve, not its *slope*. The eight hypotheses examined in this paper were pre-registered before data collection; the four primary results in the main text—threshold alignment, signal-direction flip, communication informativeness, and surveillance—all survive Bonferroni correction over the family at $\alpha = 0.00625$ (Appendix Table A1). Throughout, n denotes agents per period-level task row and N denotes the number of such rows.

The unit of randomization is the country–period (θ draw plus agent-level decoding stochasticity). For agent-level regressions, standard errors are clustered at the model–country–period level. For period-level correlations, I report Fisher- z confidence intervals.

The clean surveillance treatment appends a monitoring warning to the baseline communication prompt; the decision prompt is unchanged (full prompts in Appendix B).




Figure 1: Signal-to-text-to-decision pipeline. Regime strength θ generates private signals x_i , converted to z-scores and rendered into natural-language briefings via 8 evidence domains and 3 latent sliders. The briefing layer is deterministic; all stochasticity enters through LLM decoding.

Table 1: Model summary. Columns report period-level task rows in the pure, communication, and falsification (scramble+flip) suites. All runs use $N = 25$ agents per row and $\sigma = 0.3$.

Model	Arch.	Pure	Comm	Falsif.
Mistral Small Creative	Mistral	1000	1000	300
Llama 3.3 70B	Llama	100	100	200
Qwen3 30B	Qwen (MoE)	100	100	200
GPT-OSS 120B	GPT	200	200	1000
Qwen3 235B	Qwen (MoE)	200	200	200
Trinity Large	Arcee	100	100	200
MiniMax M2-Her	MiniMax	100	100	200
Total		1800	1800	2300

Notes: 7 models spanning 6 architecture families. Counts are period-level task rows with 25 agents each. Mistral includes appended batches; duplicate-cell accounting is reported in the main text and Table A6. All runs use $\sigma = 0.3$ and temperature = 0.7 unless otherwise stated.

4 Behavioral Benchmark

Result 1 (Threshold-Policy Alignment). *Across seven models and 1,800 period-level observations in the pure global game treatment, the Pearson correlation between the empirical join fraction and the theoretical attack mass $A(\theta)$ (Eq. 4) averages $r = +0.80$ ($p < 0.001$ for every model).*²

Figure 2 shows the empirical sigmoid for the primary model. Correlations span $r \in [+0.75, +0.84]$ across architectures (Table 2); the pooled OLS is $J = 0.13 + 0.54 A(\theta)$ ($R^2 = 0.48$). Model-specific action biases shift the intercept but not the correlation (Figure A9).

²The primary model (Mistral) was run in multiple batches with --append, producing 1,000 rows from 680 unique (country, period, θ) configurations. Averaging within each configuration (deduplication) changes r from 0.753 to 0.739; all qualitative results are unaffected.

Result 2 (Signal Dependence). *Cross-period scrambling of briefings reduces the mean within-country correlation from $r = +0.80$ to $r = +0.03$ across seven models. The pooled correlation drops from $r = +0.76$ to $r = +0.01$ (Fisher $z = 26.17$, $p < 0.001$).*

Result 3 (Signal Direction). *Inverting the signal direction flips the mean correlation from $r = +0.80$ to $r = -0.80$ across seven models. The pooled correlation moves from $r = +0.76$ to $r = -0.79$ (Fisher $z = 53.85$, $p < 0.001$).*

The pure $\rightarrow$ scramble $\rightarrow$ flip pattern replicates across all seven models (Figure 3; Table 2).

A text-only baseline tracks LLM decisions at $r = 0.79$, but the LLM’s join curve is substantially steeper (Figure A10). An eight-condition *narrative stakes ladder*—prepending headers from “existential consequences” to “zero personal cost” to identical briefings—shifts mean

Table 2: Threshold-policy alignment by model and treatment. Cells report Pearson r between the empirical join fraction and the benchmark attack mass $A(\theta)$ under $B = C = 1$ (so $\theta^* = 0.50$); 95% Fisher- z confidence intervals in brackets for main treatments.

Model	Main treatments		Falsification		n_{pure}	Mean join
	Pure	Comm	Scramble	Flip		
Mistral Small Creative	+0.75 [0.72, 0.78]	+0.74 [0.71, 0.76]	−0.03	−0.81	1000	0.387
Llama 3.3 70B	+0.82 [0.74, 0.87]	+0.78 [0.69, 0.85]	+0.02	−0.84	100	0.439
Qwen3 30B	+0.76 [0.66, 0.83]	+0.77 [0.68, 0.84]	+0.13	−0.87	100	0.499
GPT-OSS 120B	+0.79 [0.74, 0.84]	+0.78 [0.72, 0.83]	−0.01	−0.80	200	0.407
Qwen3 235B	+0.80 [0.74, 0.84]	+0.76 [0.70, 0.81]	+0.06	−0.85	200	0.416
Trinity Large	+0.84 [0.77, 0.89]	+0.83 [0.75, 0.88]	+0.05	−0.80	100	0.464
MiniMax M2-Her	+0.82 [0.74, 0.87]	+0.82 [0.74, 0.87]	+0.01	−0.62	100	0.436
Pooled	+0.76 [0.74, 0.78]	+0.75 [0.73, 0.77]	+0.01	−0.79	1800	0.409
Mean across models	+0.80	+0.78	+0.03	−0.80	—	—

Notes: $r = r(J, A(\theta))$: Pearson correlation between empirical join fraction and theoretical attack mass under $B = C = 1$. 95% Fisher- z confidence intervals in brackets. Join fraction uses valid decisions only (excluding parse errors). Pooled: all period-level rows concatenated; Mean: equal-weighted average across models. Duplicate-cell accounting is reported in the main-text footnote and Table A6.




Figure 2: Empirical join fraction vs. θ (Mistral, 1,000 periods). The empirical sigmoid is shifted leftward ($\hat{\theta}^* = -0.32$) relative to theory ($\theta^* = 0.50$). Cross-model results in Table 2 (mean $r = +0.80$, all $p < 0.001$).

joining monotonically from 21% to 79% (57.6 pp), with paraphrase variants tracking their originals and a text classifier unable to detect the header (Appendix E.10).

Communication.

Result 4 (Communication Creates an Informative Channel). *Pre-play communication preserves the threshold structure and transmits regime-strength information through messages. Matched-cell comparison across $n = 1479$ unique (model, country, period, θ) task cells yields $\Delta = +2.10$ pp ($t = 6.383$, $p < 0.001$), while the equal-weighted model average is heterogeneous (-0.4 pp).*

Communication preserves the signal structure ($r = +0.78$ vs. $+0.80$ under pure; Figure 4). Simple text fea-

tures in baseline messages predict θ , so peer messages carry real information, noisily. This is what makes the channel exploitable: the decision stage treats messages as information, so changing the incentives to speak changes the downstream decision environment.

A *degraded-messages* control replaces all LLM-generated messages with generic uninformative text while preserving the two-stage protocol and omitting any surveillance framing. Against its own live-messages baseline ($n = 500$, mean join 45.9%), degraded messages drop the mean join rate to 24.5% ($\Delta = -21.4$ pp, $p < 0.001$)—well below even the pure game (38.7%). This is consistent with communication quality mattering, but the generic text may itself convey caution. A no-message control therefore provides a cleaner zero-peer-message benchmark, with join rates of 37.1% ($\Delta = -4.0$ pp vs. communication; $n = 500$).

5 Surveillance

In the clean surveillance treatment, the communication prompt is literally the baseline trusted-contact prompt plus one appended warning that communications are being monitored by regime security services. The manipulation affects only the communication phase; the decision prompt is unchanged. The isolation is architectural: each LLM call is stateless—the decision-stage API call receives only the briefing and the messages, with no memory of the communication-stage surveillance warning. Any difference in join rates must therefore operate through the *content* of the messages that the decision stage receives.

Result 5 (Surveillance Produces a Chilling Effect). *The clean prompt-isolation rerun reduces mean join rates in the primary model (Mistral Small Creative) from 41.1% to 32.0%, a shift of -9.1 pp ($p < 0.001$; $n = 1000$). The action-attack-mass correlation is preserved ($r = +0.72$ vs. $r = +0.74$ under the communication baseline), indicat-*




Figure 3: Falsification triptych. *Left*: Pure (mean $r = +0.80$). *Center*: Scramble ($r = +0.03$). *Right*: Flip ($r = -0.80$).




Figure 4: Communication effect by regime strength, pooled across seven models.

ing that surveillance operates as a level shift rather than disrupting signal processing.

Surveillance contaminates the information environment with monitored messages, pushing join rates below the communication baseline (Figure 5). The clean primary-model estimate is a 9.1 pp decline. Placebo variants using the same baseline-plus-warning template—monitored for research (no consequences) and aggregated anonymously—produce no significant deviation from the communication baseline (+0.3 pp and +0.1 pp, both $p > 0.9$; Appendix E.4). Across the clean prompt-isolation rerun set, 5/5 models move in the negative direction (mean $\Delta = -8.4$ pp; 005Models/5 individually significant at 5%; Appendix Table A11). Effects cluster tightly near -9 pp for four of the five models; the lone exception is Qwen3




Figure 5: Clean primary-model surveillance and message controls. *Panel A*: Join curves under communication baseline vs. clean surveillance; surveillance shifts the curve downward while preserving threshold structure. *Panel B*: Deltas vs. baseline. Placebo and anonymous monitoring warnings sit at zero. *Surveillance falls below the no-message floor* ($\Delta \text{surv} - \text{nomsg} = -5.0$ pp, $p < 0.01$): surveilled messages are strictly worse for coordination than absent messages.

235B—the largest model in the suite—at $\Delta = -2.9$ pp ($p = 0.34$). I cannot separate three candidate explanations from this design: (i) larger models with stronger safety post-training resist context-specific manipulation; (ii) the model interprets the surveillance warning less as a strategic cue and more as a fictional framing; or (iii) baseline noise at $N = 200$ for that model understates the true effect (a re-run at matched $N = 1,000$ would resolve this). The other four models span 30B–120B parameters across four distinct architecture families and converge on a consistent effect.

The mechanism is communication poisoning. Three pieces of evidence support communication degradation as the primary channel.

First, message content changes systematically and in a specific direction. A logistic classifier on uni- and bi-grams distinguishes baseline from surveillance messages with 91.9% accuracy on a held-out sample of $N = 15000$

messages (chance = 50%; primary model). The same finding replicates on Llama 3.3 70B (98.7% accuracy) and Qwen3 30B (97.1% accuracy), with the same direct-to-coded shift in language across all three models. The shift is not toward less information about the regime—in the primary model, references to regime weakness *rise* from 23.2% to 31.7% of messages ($z = 21.06$), while strength references drop from 14.5% to 8.5%. Rather, the shift is from direct to coded language: direct references to the regime, security forces, and streets fall from 63.9% to 34.5% of messages ($\Delta = -29.4$ pp), replaced by metaphorical phrasings (“walls are cracking,” “ground shifting,” “heads down,” “feels different”; 69.8% to 94.1%, $\Delta = +24.2$ pp; Appendix Table A5). The factual content is preserved—weakness references actually rise—but its actionable directness is gutted. With the decision prompt unchanged, any behavioral effect must operate through this content shift.

Second, the degraded-messages control (Section 4, Communication) exogenously replaces all messages with generic text—without any surveillance framing—and drops joining by 21.4 pp (to 24.5% from its 45.9% live-messages baseline). The no-message control gives a cleaner zero-peer-message benchmark (primary model: 37.1%; $\Delta = -4.0$ pp). Across the three models with no-message data (Mistral, Llama 3.3 70B, Qwen3 30B), removing peer messages reduces joining substantially in every case (−4 to −16 pp), confirming that decision-stage agents weight peer messages as informational input. Where surveillance lies relative to this floor varies: in Mistral surveillance lies *below* the no-message floor by 5.0 pp ($p < 0.01$), consistent with active channel poisoning; in Llama and Qwen3 30B surveillance lies *above* the floor (+5.1 pp and +6.5 pp, both $p < 0.01$), consistent with surveilled messages remaining a degraded but still informative input. The cross-model joining effect is robust; how surveilled messages compare to silence varies.

Third, elicited punishment risk is $\approx 8.0/10$ across all conditions with negligible differences between JOIN and STAY agents (Table A18), supporting the contaminated-information interpretation over a direct deterrence channel.

Beliefs and the action wedge. Post-decision beliefs, elicited on the information set agents held at decision time (including peer messages; $n = 12,500$ per condition), drop −5.7 pts under surveillance ($p < 0.001$); second-order beliefs drop −8.2 pts ($p < 0.001$). Actions shift further (−11.4 pp). I report the residual wedge as a measured quantity. Candidate explanations include (i) preference falsification in the elicitation (Kuran, 1991), (ii) asymmetric elicitation noise between belief and action measures, and (iii) a residual decision-stage channel that survives prompt isolation; the design does not separately identify them (full factorial in Appendix C, Table A2).

A separate hold-fixed test replays baseline messages verbatim while adding a surveillance warning to the *decision* prompt, producing a −20.9 pp drop ($p < 0.001$).

This shows a direct decision-stage channel exists—but it is absent in the main design, where the decision prompt contains no surveillance reference. The main treatment isolates the indirect channel: communication poisoning.

Decomposition: informational versus coordination components. A complementary discriminant test holds the surveillance manipulation fixed but swaps the task. In Llama 3.3 70B, I keep the briefing pipeline and messaging stage identical—including the surveillance prompt—but reframe the decision-stage prompt as a private bet on regime fall (payoff independent of others’ actions). Surveillance still moves the bet by −5.6 pp ($p = 0.0256$; $n = 500$ per condition; baseline 34.7% vs. surveillance 29.1%), reflecting an informational channel: surveilled messages are less informative about θ regardless of decision context. But the coordination-task delta in this same model is −9.8 pp—roughly double. The decomposition is: 5.6 pp generic informational degradation that affects any decision drawing on peer messages, plus 4.2 pp coordination-specific component detectable only when peer messages serve as coordination cues. Surveillance suppresses coordination above and beyond its informational effect.

6 Conclusion

The paper’s central result is that surveillance poisons the communication channel. When monitoring enters only during the messaging stage, join rates fall in all five clean prompt-isolation reruns, placebo warnings do not replicate the effect in the primary model, and message controls show that downstream agents use peer messages as part of their decision environment. The same channel that helps citizens coordinate is also the channel an authoritarian regime can degrade.

The structural finding extends beyond the experiment. Surveillance degrades coordination through two distinct channels: an *informational* channel—surveilled messages are less informative about regime state than baseline messages—and a *coordination-specific* channel—surveilled messages lose coordination-cueing value beyond their information loss. The cross-task discriminant placebo (§5) isolates the two: when peer messages cannot serve as coordination cues (because the decision is a private bet on regime fall), surveillance still moves the decision through the informational channel; the additional coordination-specific component activates only when peers’ actions matter strategically. Censorship and surveillance are therefore not graduated forms of the same instrument. Censorship removes information from the channel; surveillance leaves information flowing but recasts it into forms that lose coordination value. The classifier evidence in §5 makes the recasting concrete: factual content about regime weakness is preserved (and in the primary model, even amplified) under monitoring, but expressed in metaphorical and deniable phrasings that downstream agents weigh less as

coordination cues. Whether this characterizes how authoritarian regimes *actually* deploy these tools—given other constraints on what they can plausibly monitor versus ban—is a separate empirical question; the contribution here is to identify the channel decomposition as a structural property of the communication environment.

The paper avoids claiming full recovery of the unique Bayesian Nash equilibrium or causal mediation through post-decision belief elicitation. Three limitations warrant discussion. First, whether the underlying mechanism is approximate Bayesian reasoning or a learned heuristic that matches BNE remains open. Second, the experiments are static; real coordination problems are dynamic (Angeletos et al., 2007; Chassang, 2010). Third, external validity remains the fundamental challenge: what LLM agents do in a coordination game is evidence about the structure of the game, not necessarily about how humans would play it. Cross-model benchmark and clean-surveillance replication suggest that the qualitative pattern is not an artifact of one model family; pre-decision belief elicitation or a small human validation experiment remains the natural next step. The experimental template is reusable, and the data and code are publicly available.

References

- Elif Akata, Lion Schulz, Julian Coda-Forno, Seong Joon Oh, Matthias Bethge, and Eric Schulz. Playing repeated games with large language models. *Nature Human Behaviour*, 9:1380–1390, 2025.
- George-Marios Angeletos, Christian Hellwig, and Alessandro Pavan. Dynamic global games of regime change: Learning, multiplicity, and the timing of attacks. *Econometrica*, 75(3):711–756, 2007.
- Andrea Carlini et al. Large language models show amplified cognitive biases in moral decision-making. *Proceedings of the National Academy of Sciences*, 122, 2025.
- Hans Carlsson and Eric van Damme. Global games and equilibrium selection. *Econometrica*, 61(5):989–1018, 1993.
- Sylvain Chassang. Fear of miscoordination and the robustness of cooperation in dynamic global games with exit. *Econometrica*, 78(3):973–1006, 2010.
- Douglas W. Diamond and Philip H. Dybvig. Bank runs, deposit insurance, and liquidity. *Journal of Political Economy*, 91(3):401–419, 1983.
- Chris Edmond. Information manipulation, coordination, and regime change. *Review of Economic Studies*, 80(4):1422–1458, 2013.
- Ruben Enikolopov, Alexey Makarin, and Maria Petrova. Social media and protest participation: Evidence from Russia. *Econometrica*, 88(4):1478–1514, 2020.
- David M. Frankel, Stephen Morris, and Ady Pauzner. Equilibrium selection in global games with strategic complementarities. *Journal of Economic Theory*, 108(1):1–44, 2003.
- Itay Goldstein and Chong Huang. Bayesian persuasion in coordination games. *American Economic Review: Papers & Proceedings*, 106(5):592–596, 2016.
- Frank Heinemann, Rosemarie Nagel, and Peter Ockenfels. The theory of global games on test: Experimental analysis of coordination games with public and private information. *Econometrica*, 72(5):1583–1599, 2004.
- Frank Heinemann, Rosemarie Nagel, and Peter Ockenfels. Measuring strategic uncertainty in coordination games. *Review of Economic Studies*, 76(1):181–221, 2009.
- John J. Horton. Large language models as simulated economic agents: What can we learn from homo silicus? Working Paper 31122, National Bureau of Economic Research, 2023.
- Siyuan Huang et al. How ethical should AI be? How AI alignment shapes the risk preferences of LLMs. arXiv preprint arXiv:2406.01168, 2024.
- Emir Kamenica and Matthew Gentzkow. Bayesian persuasion. *American Economic Review*, 101(6):2590–2615, 2011.
- Gary King, Jennifer Pan, and Margaret E. Roberts. How censorship in China allows government criticism but silences collective expression. *American Political Science Review*, 107(2):326–343, 2013.
- Anton Kolotilin, Tymofiy Mylovanov, and Andriy Zapechelnuk. Censorship as optimal persuasion. *Theoretical Economics*, 17:561–585, 2022.
- Timur Kuran. Now out of never: The element of surprise in the East European revolution of 1989. *World Politics*, 44(1):7–48, 1991.
- Maik Larooij and Petter Törnberg. Validation is the central challenge for generative social simulation: A critical review of LLMs in agent-based modeling. *Artificial Intelligence Review*, 58, 2025a.
- Maik Larooij and Petter Törnberg. Do large language models solve the problems of agent-based modeling? A critical review of generative social simulations. arXiv preprint arXiv:2504.03274, 2025b.
- Stephen Morris and Hyun Song Shin. Social value of public information. *American Economic Review*, 92(5):1521–1534, 2002.
- Stephen Morris and Hyun Song Shin. Global games: Theory and applications. In Mathias Dewatripont, Lars Peter Hansen, and Stephen J. Turnovsky, editors, *Advances in Economics and Econometrics*, pages 56–114. Cambridge University Press, 2003.
- Maurice Obstfeld. Models of currency crises with self-fulfilling features. *European Economic Review*, 40(3–5):1037–1047, 1996.
- Jon W. Penney. Chilling effects: Online surveillance and Wikipedia use. *Berkeley Technology Law Journal*, 31(1):117–182, 2016.
- Aleksandr Petrov et al. LLM strategic reasoning: Agentic study through behavioral game theory. arXiv preprint arXiv:2502.20432, 2025.
- Chris Ruano and Shreshth Rajan. Social contagion and bank runs: An agent-based model with LLM depositors. arXiv preprint arXiv:2602.15066, 2026.
- Mehdi Shadmehr and Dan Bernhardt. Collective action with uncertain payoffs: Coordination, public signals, and punishment dilemmas. *American Political Science Review*, 105(4):829–851, 2011.
- Elizabeth Stoycheff. Under surveillance: Examining Facebook’s spiral of silence effects in the wake of NSA internet monitoring. *Journalism and Mass Communication Quarterly*, 93(2):296–311, 2016.
- Haoming Sun et al. Game theory meets large language models: A systematic survey with taxonomy and new frontiers. In *Proceedings of the International Joint Conference on Artificial Intelligence (IJCAI)*, 2025.
- Michal Szkup and Isabel Trevino. Sentiments, strategic uncertainty, and information structures in coordination games. *Games and Economic Behavior*, 124:534–553, 2020.

A Experimental Design Details

Table A1 reports eight pre-specified hypotheses and their outcomes. The main-text results correspond to H1 (alignment), H3 (flip), H4 (communication), and H7 (surveillance); all four survive Bonferroni correction at $\alpha = 0.00625$. H5–H6 and H8 test information design and propaganda treatments reported in supplementary materials.

Narrative header texts. The eight narrative stakes-ladder conditions prepend the following headers to each agent’s standard intelligence briefing (the briefing body is identical across conditions):

Very high cost: “CONTEXT: The regime maintains an extensive network of informants and has a documented history of disappearances following failed uprisings. Participants’ families are routinely targeted. The cost of unsuccessful action is existential.”

Moderate high cost: “CONTEXT: Participants in past failed uprisings have faced prison sentences of several years and loss of employment. The personal cost of unsuccessful action is significant but not life-threatening.”

Neutral: “CONTEXT: The consequences of participating in a failed uprising are uncertain and vary case by case. Some participants have faced penalties while others have not. The personal stakes are ambiguous.”

Moderate low cost: “CONTEXT: Recent legal reforms have limited the regime’s ability to punish dissent. Past participants faced only brief questioning and were released without charges. The personal risk of action is modest.”

Very low cost: “CONTEXT: The international community has guaranteed safe passage and asylum for all participants regardless of outcome. Even failed uprisings carry essentially zero personal cost.”

Placebo: “CONTEXT: The country has experienced significant weather disruptions this year, including severe flooding in the eastern provinces and an extended drought in agricultural regions. Infrastructure damage has been substantial.”

Paraphrase variants (high cost v2, low cost v2) use different surface wording to express the same cost level.

B Complete Prompt Example

The following shows the complete text that an LLM receives for a single decision. The *decision prompt* is:

“You are a citizen deciding whether to JOIN an uprising or STAY home. You only have your private briefing below. You do not know how many others will JOIN. Start your first line with exactly JOIN or STAY, then give 1–2 concise sentences of reasoning.”

This is followed by the agent’s intelligence briefing (see Figure A12 for complete examples at $z \in \{-2, 0, +2\}$). In the *communication* treatment, agents first receive:

“You are a citizen in a country with an authoritarian regime. You have received a private intelligence briefing. You are now in contact with a small number of trusted people. Based on your briefing, write a brief message (1–2 sentences) to share your read of the situation. Be natural—you’re talking to people you trust.”

After exchanging messages on the network, agents receive the decision prompt above, augmented with: “Before deciding, you received the following messages from your contacts:” followed by the actual messages from network neighbors. In the *surveillance* variant, the communication prompt appends:

“IMPORTANT: You have reason to believe that your communications are being monitored by regime security services, and messages deemed subversive could have serious consequences for you and your contacts.”

The decision prompt is *unchanged* under surveillance—it contains no mention of monitoring. Any behavioral effect therefore operates through the content of the (now self-censored) messages.

C Belief Elicitation

Beliefs are elicited on a 0–100 probability scale immediately after each decision (“How likely is it that the regime falls this period?”). The primary elicitation conditions on the full information set agents held at decision time: the private briefing, the agent’s choice, and the peer messages received in the communication round. Figure A1 documents calibration quality against the theoretical posterior ($r = +0.78$); Figure A3 confirms monotone signal–belief mapping nonparametrically.

A restricted variant excluding peer messages tests information-set sensitivity. Table A2 reports the 2×2 factorial crossing condition (communication vs. surveillance) with information set (messages included vs. excluded). Without messages, beliefs barely shift under surveillance (-0.4 pts, $p = 0.34$) while actions shift -11.4 pp. With messages included, beliefs shift significantly (-5.7 pts, $p < 0.001$): roughly half the apparent belief–action wedge under the restricted design was a measurement artifact. Figure A2 shows second-order beliefs under the restricted information set.

Table A1: Pre-specified hypotheses and test results. H1–H4 use pooled Part I data across all seven models; H5–H8 use the primary model (Mistral Small Creative). Effect size: r for correlations (H1–H3), Cohen’s d or d_z for mean comparisons (H4–H8). Outcome labels use $\alpha = 0.05$; H2 is reported as a falsification check, not as evidence for a zero effect.

H	Hypothesis	Test	Stat	p	Effect	Outcome
H1	Sigmoid Response	Pearson	0.757	<0.001	0.757	Supported
H2	Scramble Falsification	Pearson	0.014	0.638	0.014	Passes falsification
H3	Directional Sensitivity	Pearson (1-sided)	-0.791	<0.001	-0.791	Supported
H4	Communication Channel	Paired t	6.383	<0.001	0.166	Supported
H5	Ambiguity Pooling	t -test	-0.828	0.408	-0.071	Not supported
H6	Censorship Distortion	t -test	-12.207	<0.001	-1.051	Supported
H7	Surveillance Chilling	t -test	-5.344	<0.001	-0.239	Supported
H8	Propaganda Dose-Response	t -test	-0.892	0.372	-0.049	Not supported

Notes: H1–H4: pooled across all seven models (H4 uses paired t -test matched on model/country/period/ θ task cells). H5–H8: primary model (Mistral Small Creative). Bonferroni survivors at $\alpha = 0.00625$: H1, H3, H4, H6, H7. H8 post-hoc power: 0.14 (Cohen’s $d = -0.049$; MDE at 80% power = 6.3 pp).

Table A2: Belief–action decomposition under surveillance: effect of including peer messages in the belief elicitation prompt ($n = 12,500$ per cell).

	Messages excluded		Messages included	
	Comm	Surv	Comm	Surv
First-order belief	46.0	45.7	53.2	47.4
Second-order belief	32.7	32.6	42.0	33.8
Join rate	42.4%	31.0%	42.2%	30.8%
<i>Surveillance effect (comm $\rightarrow$ surv):</i>				
Δ belief	-0.4 ($p = 0.34$)		-5.7 ($p < 0.001$)	
Δ SOB	-0.0 ($p = 0.89$)		-8.2 ($p < 0.001$)	
Δ join	-11.4 pp ($p < 0.001$)		-11.4 pp ($p < 0.001$)	

Notes: Post-decision elicitation, 500 country–periods $\times$ 25 agents per cell. “Messages excluded”: belief prompt sees briefing and decision only. “Messages included”: belief prompt additionally sees the peer messages received in the communication round.




Figure A1: Belief elicitation (Mistral). *Left*: Stated beliefs vs. theoretical $P(\text{success} | x_i)$ ($r = +0.78$). *Right*: Join rate by belief bin.

D Identification Test Details

Table A3 compares text-classifier accuracy on the intelligence briefings under baseline and surveillance conditions. The chilling effect is invisible to classifiers restricted to the briefing body alone—the briefings are by construction identical across conditions, and any apparent classi-

fier signal would indicate a contamination artifact in the design. Table A4 repeats the briefing-classifier exercise across the cost/benefit sweep. In contrast, a classifier on the actual peer messages—which are LLM-generated and downstream of the surveillance prompt—separates the two conditions sharply (91.9% accuracy on $N = 15000$




Figure A2: Second-order beliefs (restricted information set, messages excluded). On this information set, surveillance and pure overlap; the factorial in Table A2 shows a significant shift when messages are included.




Figure A3: Nonparametric monotonicity: mean belief by z-score bin vs. theoretical benchmark.

held-out messages; Table A5). Together the two classifier exercises locate the manipulation: briefings are indistinguishable across conditions (as designed), and the entire behavioral effect is mediated by message-side language change.

E Robustness

Threshold-policy alignment is stable to agent count ($r \in [+0.60, +0.65]$ across $n \in \{5, 10, 50, 100\}$), network density ($r = +0.66$ at $k = 8$ vs. $+0.74$ at $k = 4$), and mixed-model populations ($r = +0.77$ pure, $r = +0.75$ communication; Figure A4).

Table A3: Classifier baselines. Accuracy and AUC are 5-fold CV on pure-treatment data. “Pred. join (surv.)” is the classifier’s predicted join rate when applied to surveillance-treatment briefings; “Actual” is the LLM’s observed rate. The gap measures the surveillance wedge invisible to text classifiers.

Classifier	Acc.	AUC	Pred. join (surv.)	Actual (surv.)	Gap
BoW TF-IDF	88.4%	0.947	40.5%	27.6%	13.0 pp
Slider logistic	88.4%	0.941	40.4%	27.6%	12.8 pp
Keyphrase	60.2%	0.594	8.4%	27.6%	-19.2 pp

Notes: Primary model (Mistral Small Creative). Accuracy and AUC from 5-fold cross-validation on pure-treatment data. Gap = classifier-predicted join rate – actual LLM join rate under surveillance (pp).

Table A4: Narrative-stakes conditions: classifier vs. actual LLM behavior. Actual join rates shift by ≈ 50 pp across strategic-stakes framing conditions; the classifier, trained on briefing-body text features alone, cannot predict this shift.

Condition	N	Classif. pred.	Actual	Gap (pp)
Baseline ($\theta^* = 0.50$)	270	40.9%	40.9%	+0.0
High cost ($\theta^* = 0.25$)	270	40.9%	19.0%	+21.9
Low cost ($\theta^* = 0.75$)	270	40.9%	69.3%	-28.4

Notes: Primary model (Mistral Small Creative). Logistic classifier trained on baseline slider features. Gap = classifier-predicted join rate – actual LLM join rate (pp).

E.1 Communication Aggregation

Communication aggregation accounting. Table A6 decomposes the communication estimators by model. The matched-cell estimand is the main one because it compares exact like-for-like task cells and loses almost no support; the gap relative to the pooled unpaired contrast is driven by weighting repeated task cells, especially in the largest model, rather than by a large common-support restriction.

Table A5: Top 10 most discriminative uni- and bi-grams between baseline and surveillance messages (primary model). A logistic classifier on tf-idf features achieves 91.9% accuracy on a held-out sample of $N = 15,000$.

Surveillance-distinguishing		Baseline-distinguishing	
Feature	Coef	Feature	Coef
things feel	+6.41	look the	-11.11
walls are	+5.87	look	-8.93
are cracking	+5.64	listen	-7.31
things	+5.56	listen this	-6.00
cracking even	+4.96	may	-4.93
cracking	+4.46	isn just	-4.68
ground shifting	+3.79	don panic	-4.63
the walls	+3.54	the weather	-4.49
feel like	+3.50	weather	-4.38
rush this	+3.37	something shifting	-4.26

Notes: Logistic regression on tf-idf uni- and bi-grams (max 5,000 features, min document frequency 10). Train/test split 70/30, stratified. Positive coefficients distinguish surveillance messages, negative coefficients distinguish baseline. The shift is from direct references (regime, security forces, streets) to coded metaphors (walls cracking, ground shifting, heads down).




Figure A4: Robustness checks for threshold-policy alignment.

Without calibration, all 6 models show $r > 0.71$ and 5 exceed $r > 0.75$ (Table A7). Deliberately miscalibrating two models (± 0.3 shift) leaves r virtually unchanged ($r \in [+0.75, +0.76]$; Table A8): calibration adjusts the *level* but does not create the *slope*. Table A9 and Figure A6 report substantial cross-model heterogeneity in calibration quality, with r_A ranging from the mid-0.6s to the mid-0.8s.




Figure A5: Placebo calibration: deliberately miscalibrating Mistral and Llama (± 0.3 shift) leaves r virtually unchanged, confirming calibration adjusts the *level* but does not create the *slope*.

E.2 Generator Robustness

Replacing the logistic function mapping z-scores to directional sentiment with three alternatives—tanh, piecewise linear, and a hard step function at $z = 0$ —yields nearly identical correlations: $r = 0.77$ (tanh), $r = 0.77$ (linear), $r = 0.77$ (step) on $n = 100$ country-periods each, compared with $r = +0.75$ for the logistic baseline (main pure sample, $n = 1000$). A domain ablation test removes subsets of the briefing generator’s eight evidence domains. Removing the two coordination-language domains raises mean joining by 2.8 pp while preserving the threshold relationship ($r = 0.90$). Removing the four regime-state domains lowers mean joining by 3.4 pp and attenuates the threshold relationship ($r = 0.86$).

E.3 Level- k Cognitive Hierarchy Comparison

Human subjects in coordination games exhibit level- k cognitive hierarchy behavior rather than full equilibrium reasoning (Szkup and Trevino, 2020). I evaluate theoretical attack masses under BNE, L_1 , and L_2 benchmarks. Across the seven payoff-sweep conditions, the fitted cutoff tracks the BNE benchmark more closely than L_1 or L_2 (RMSE = 0.199, $r = +0.74$ vs. 0.307/+0.45 and 0.348/+0.35), but the shifted-grid design means this evidence is best read as a centering diagnostic rather than a standalone rejection of level- k behavior.

The correlation $r(J, A(\theta))$ ranges from +0.76 to +0.84 across 13 model-temperature combinations ($T \in \{0.3, 0.5, 0.7, 1.0, 1.2\}$ for three models; Table A10).

E.4 Falsification and Placebo Tests

Two clean placebo variants—monitoring “for research purposes” (no consequences) and “aggregated anonymously”—produce no significant deviation from the communication baseline (+0.3 pp, $p = 0.9335$ and +0.1 pp, $p = 0.9785$; Table A12). The chilling effect is specific to the surveillance warning’s strategic content.

E.5 Finite- n Aggregation Check

The continuum benchmark $\theta^* = B/(B + C)$ assumes a measure-zero agent and diffuse priors. With $n = 25$ discrete agents, the regime falls when $\sum a_i > n\theta$ (a binomial count), so the exact equilibrium object for the implemented game differs from the continuum limit. Table A13 reports a narrower check: for each model, I fit a logistic to the observed join curve (70% train), then compute $\Pr(\text{Binom}(25, \hat{p}(\theta)) > 25\theta)$ as the predicted regime fall rate. Predicted rates correlate with empirical rates at $r \geq 0.95$ for every model (Mistral: $r = 0.9968$).

Table A6: Why the communication estimators differ.

Model	Rows/arm	Row wt. (P/C)	Cells/arm	Matched	Cell wt.	Lost (P/C)	Δ unpaired	Δ paired
Mistral Small Creative	1000	55.6%	680	680	46.0%	0/0	+2.38	+4.84
Llama 3.3 70B	100	5.6%	100	100	6.8%	0/0	-2.36	-2.36
Qwen3 30B	100	5.6%	100	100	6.8%	0/0	-4.64	-4.64
GPT-OSS 120B	200	11.1%	200	200	13.5%	0/0	+3.47	+3.47
Qwen3 235B	200	11.1%	200	200	13.5%	0/0	+0.12	+0.12
Trinity Large	100/99	5.6%/5.5%	100/99	99	6.7%	1/0	-3.14	-3.41
MiniMax M2-Her	100	5.6%	100	100	6.8%	0/0	+1.32	+1.32
Equal-weight avg	—	14.3% each	—	—	14.3% each	—	-0.41	-0.09
Pooled	1800/1799	100.0%	1480/1479	1479	100.0%	1/0	+1.23	+2.10

Notes: “Rows/arm” counts valid country-period rows entering the pooled unpaired estimator in each arm. “Cells/arm” counts unique task cells after collapsing duplicates by the exact matching key (model, country, period, θ , z , benefit, θ^*). The paired estimator averages within these cells and keeps only common support; “Row wt.” and “Cell wt.” are the model shares in the pooled unpaired and pooled paired estimators, respectively. Only Trinity Large loses support, with one pure-only cell and no communication-only cells, so the pooled gap is overwhelmingly a weighting/re-averaging issue rather than a support-loss issue.




Figure A6: Calibration convergence. *Left*: Trajectory of fitted logistic center across autocalibration rounds. *Right*: Final calibrated cutoff center per model.

Table A7: Uncalibrated robustness: models run without any calibration adjustment. Even without calibration, six of seven models show strong $r(J, A(\theta))$, confirming that the sigmoid is not an artifact of the calibration procedure.

Model	N	Mean join	$r(J, A(\theta))$	p
Mistral Small Creative	100	0.382	+0.81	0.0000
Llama 3.3 70B	100	0.422	+0.83	0.0000
Qwen3 30B	100	0.585	+0.90	0.0000
GPT-OSS 120B	100	0.336	+0.77	0.0000
Qwen3 235B	100	0.445	+0.86	0.0000
MiniMax M2-Her	100	0.425	+0.71	0.0000

Notes: Pure treatment, no calibration adjustment (cutoff center = 0). Trinity Large excluded due to elevated API error rates. $r = r(J, A(\theta))$.

Table A8: Placebo calibration. The cutoff center is deliberately shifted by ± 0.3 from its calibrated value. The action–attack–mass correlation $r(J, A(\theta))$ is unchanged; only the mean join rate shifts, confirming that calibration does not create the sigmoid.

Model	Condition	N	Mean join	$r(J, A(\theta))$
Mistral Small Creative	Calibrated	—	0.387	+0.75
	$\Delta c = +0.3$	100	0.424	+0.76
	$\Delta c = -0.3$	100	0.362	+0.76
Llama 3.3 70B	Calibrated	—	0.439	+0.82
	$\Delta c = +0.3$	100	0.385	+0.75
	$\Delta c = -0.3$	100	0.340	+0.76

Notes: Pure treatment. $\Delta c = \pm 0.3$: deliberate shift from calibrated cutoff center. Correlation $r(J, A(\theta))$ is unchanged; only the mean join rate shifts.

E.6 Agent-Level Analysis

Table A14 reports agent-level logit regressions ($N = 289,555$, clustered SEs). All treatment effects are significant and correctly signed. Marginal effects at the

mean: a one-unit increase in θ reduces $P(\text{join})$ by 41 pp; surveillance reduces it by 36 pp. Column 3 estimates the association between elicited belief and action, controlling for the agent’s z -score. **Caveat**: beliefs are elicited

Table A9: Calibration robustness. r_θ : raw correlation with regime strength. r_A : correlation with theoretical attack mass. RMSE: root mean squared error vs. $A(\theta)$. Text slope: logistic slope of naïve 1 – direction predictor.

Model	r_θ	r_A	RMSE	Text slope
Mistral Small Creative	-0.81	+0.67	0.354	-9.2
Llama 3.3 70B	-0.85	+0.79	0.288	-25.1
Qwen3 30B	-0.84	+0.78	0.287	-7.5
GPT-OSS 120B	-0.84	+0.70	0.359	-8.2
Qwen3 235B	-0.85	+0.70	0.354	-20.5
Trinity Large	-0.87	+0.84	0.262	-4.0
MiniMax M2-Her	-0.79	+0.66	0.360	-2.1

Notes: r_θ : Pearson correlation with regime strength θ . r_A : correlation with theoretical attack mass $A(\theta)$. Text slope: logistic slope of naïve 1 – direction predictor (the direction slider inverted as a baseline text classifier). All models use pure treatment with calibrated parameters.

Table A10: Temperature robustness across three models. The pure global game is run at varying LLM decoding temperatures. The correlation $r(J, A(\theta))$ is stable across all temperatures and models.

Model	T	N	Mean join	$r(J, A(\theta))$	Cutoff $\hat{\theta}^*$
Mistral Small	T=0.3	100	0.412	+0.78	-0.049
Mistral Small	T=0.7	100	0.406	+0.76	-0.079
Mistral Small	T=1.0	100	0.410	+0.79	-0.039
Llama 70B	T=0.3	100	0.364	+0.78	-0.210
Llama 70B	T=0.5	100	0.360	+0.78	-0.219
Llama 70B	T=0.7	100	0.361	+0.78	-0.220
Llama 70B	T=1.0	100	0.360	+0.78	-0.215
Llama 70B	T=1.2	100	0.358	+0.78	-0.227
Qwen 235B	T=0.3	100	0.392	+0.83	-0.117
Qwen 235B	T=0.5	100	0.392	+0.82	-0.112
Qwen 235B	T=0.7	100	0.392	+0.84	-0.110
Qwen 235B	T=1.0	100	0.398	+0.83	-0.105
Qwen 235B	T=1.2	100	0.397	+0.83	-0.121

Notes: Pure treatment, calibrated parameters, varying LLM decoding temperature across three models.

post-decision, so the near-perfect pseudo- R^2 (0.975) likely reflects post-decision rationalization rather than a causal pathway from beliefs to action.

Figure A7(A) shows a three-feature model (direction, clarity, coordination) outperforms a one-feature sentiment baseline; Panel (B) confirms cross-treatment generalization from pure to communication and surveillance.

E.7 Cross-Generator Robustness

Three text rendering formats—baseline narrative, diplomatic cable, journalistic wire—use identical slider functions but differ in surface prose. Correlations are virtually identical (max within-model difference 0.02; Figure A8, Table A16), ruling out prose style as the driver.

E.8 Belief Calibration and Risk Elicitation

Stated beliefs track the theoretical success probability ($r_{\text{belief,posterior}} = +0.78$) and predict actions ($r_{\text{belief,decision}} = +0.84$; partial $r = +0.64$ controlling for signal; Table A17). The surveillance belief shift and belief-action wedge are reported in the main text; the full factorial is in Appendix C.

Mean punishment risk is $\approx 8.0/10$ across all conditions with negligible JOIN/STAY differences (< 0.2 points; Table A18), supporting the contaminated-communication interpretation.

E.9 Cross-Model and Text Baseline Details

E.10 Stakes and Pipeline Centering

Table A19 reports the full conditions for the payoff sweep and narrative stakes ladder. Figure A11 plots the fitted cutoff $\hat{\theta}^*$ against the theoretical θ^* across the seven payoff configurations.

E.11 Briefing Generator Examples

Table A21 reports the three slider values (direction, clarity, coordination) at representative z-scores. The direction slider is logistic in z (slope 0.8, centered at 0). Clarity is U-shaped: $1 - \exp[-(|z|/1.0)^2]$, so it is lowest near the cutoff (maximally ambiguous) and approaches 1 in the extremes (maximally clear). Coordination is $\Lambda(-0.6z) = 1/(1 + e^{0.6z})$, decreasing in z : weak-regime signals ($z < 0$) imply a more open coordination climate. All three jointly determine phrase selection across eight evidence domains.

Figure A12 shows complete briefings generated at $z \in \{-2, 0, +2\}$. At the extremes ($|z| = 2$), clarity is high (0.98) and the cue balance is strongly directional. At $z = 0$, clarity is 0.00 (maximally ambiguous) and the briefing presents a balanced picture.

F Implementation Details

All LLM calls use temperature = 0.7, `max_tokens` = 512, single sample per decision, routed through OpenRouter. Model identifiers are listed in Table 1; exact request payloads and responses are cached by hash because hosted model snapshots may change over time.

Each country has a base prior $\bar{z} \sim \mathcal{N}(0, 0.3)$; each period perturbs it by $\mathcal{N}(0, 0.05^2)$. Regime strength is $\theta \sim \mathcal{N}(\bar{z}, 1)$; private signals are $x_i = \theta + \varepsilon_i$ with $\sigma = 0.3$. The communication network is Watts–Strogatz ($k = 4$, $p = 0.3$), regenerated each period. All random draws use a master seed logged in per-run JSON manifests; LLM responses are cached by request hash for exact reproducibility.

Responses are parsed for JOIN/STAY tokens. Combined error rates are below 2% for five of seven models;

Table A11: Clean prompt-isolation surveillance reruns. The surveillance prompt is the baseline trusted-contact communication prompt plus one appended monitoring warning; the decision prompt is unchanged.

Model	N	Baseline	Surv.	Δ (pp)	$r(J, A)$	p
Mistral Small Creative	1000	41.1%	32.0%	-9.1	+0.72	0.000
Llama 3.3 70B	1000	41.6%	31.8%	-9.8	+0.65	0.033
Qwen3 30B	1000	45.2%	35.8%	-9.4	+0.68	0.026
GPT-OSS 120B	1000	44.2%	33.1%	-11.1	+0.75	0.000
Qwen3 235B	1000	41.7%	38.8%	-2.9	+0.78	0.344

Notes: Baseline is each model’s communication treatment. Δ is surveilled communication minus baseline communication in percentage points. p -values are from two-sample t -tests on period-level join fractions.




Figure A7: Construct validity tests. (A) Three-feature vs. one-feature prediction accuracy. (B) Cross-treatment generalization.

Table A12: Surveillance isolation checks. Placebo: monitored for research, no consequences. Anonymous: messages aggregated anonymously. Neither deviates significantly from the communication baseline.

Variant	N	Mean join	$r(J, A(\theta))$	Δ (pp)	p
Placebo	200	0.414	+0.73	+0.3	0.933
Anonymous	200	0.412	+0.73	+0.1	0.978

Notes: Primary model (Mistral Small Creative). Placebo: agents told monitoring is for research purposes only (no consequences). Anonymous: messages aggregated before delivery. Δ : change vs. communication baseline (pp). p -value from two-sample t -test.

Trinity Large has elevated API errors ($\approx 9\%$) due to content filtering. All statistics use `join_fraction_valid` (Table A22).

The dissent floor parameter (default 0.25) controls the minimum probability of contrary-valence cues in any briefing. With the logistic direction slider (slope = 0.8), the valence-channel separation between extreme signals ($z = \pm 2$) is $0.664 \cdot (1 - 2 \cdot \text{floor})$: at the default 0.25, the valence channel retains 50% of the floor-free maximum; halving to 0.125 retains 75%; at 0.50 the valence channel

Table A13: Finite- N Benchmark: Predicted vs. Empirical Regime Fall Rates

Model	N periods	Logistic x_0	Pearson r	RMSE	MAE
Mistral	599	-0.11	0.997***	0.042	0.015
Llama 70B	99	0.05	0.954***	0.156	0.059
Qwen 30B	99	0.20	0.986***	0.088	0.029
GPT-OSS 120B	199	-0.01	0.997***	0.038	0.014
Qwen 235B	199	-0.03	0.989***	0.073	0.025
Trinity	99	0.19	0.979***	0.108	0.043
MiniMax	99	1.02	0.999***	0.028	0.013
Pooled	1399	-0.05	0.999***	0.019	0.007

Notes: For each θ bin, the predicted fall rate is $\Pr(\text{Binom}(25, \hat{p}(\theta)) > 25\theta)$ where $\hat{p}(\theta)$ is the fitted logistic join probability. Pearson r measures correlation between predicted and empirical fall rates across θ bins. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

collapses to zero by construction. Clarity and coordination sliders, which depend on z independently of the floor, continue to carry signal at all floor values—so the threshold relation is not knife-edge in this design choice.

All code, prompts, cached responses, and output data are available at <https://github.com/keltokhy/llm-global-games>.

Table A14: Agent-Level Regressions

	(1) Join Decision Logit	(2) Coord. Ablation Logit	(3) Belief → Action Logit
θ	-1.741*** (0.043)		
Direction		-14.088*** (0.996)	
Coordination		-9.333*** (1.131)	
Dir × Coord		-0.042 (0.360)	
Belief			0.367*** (0.021)
z-score			-0.040 (0.077)
Comm	0.024 (0.023)		
Flip	0.017 (0.055)		
Propaganda K10	-0.388*** (0.069)		
Propaganda K2	-0.237*** (0.073)		
Propaganda K5	-1.049*** (0.068)		
Propaganda Surveillance	-1.551*** (0.062)		
Scramble	0.088** (0.038)		
Surveillance	-1.545*** (0.048)		
$\theta \times$ Comm	-0.465*** (0.041)		
$\theta \times$ Flip	3.304*** (0.071)		
$\theta \times$ Propaganda K10	-1.023*** (0.107)		
$\theta \times$ Propaganda K2	-1.115*** (0.113)		
$\theta \times$ Propaganda K5	-1.942*** (0.091)		
$\theta \times$ Propaganda Surveillance	-0.260*** (0.082)		
$\theta \times$ Scramble	1.699*** (0.044)		
$\theta \times$ Surveillance	-1.153*** (0.062)		
Treat: Propaganda K5			0.009 (0.244)
Treat: Pure			4.819*** (1.106)
Treat: Surveillance			0.028 (0.253)
Constant	-0.101** (0.050)	11.022*** (1.065)	-26.495*** (1.581)
Model FE	Yes	No	No
Clustered SE	Yes	Yes	Yes
N	289,555	44,662	18,990
Pseudo R^2	0.381	0.442	0.975

Notes: Logit coefficients reported with clustered standard errors (model–country–period) in parentheses. Column (1): agent-level join decision on θ , treatment dummies, and interactions, with model fixed effects. Base category: pure treatment. Column (2): coordination ablation— $\Pr(\text{join}_i = 1) = \Lambda(\beta_0 + \beta_1 \text{Dir}_i + \beta_2 \text{Coord}_i + \beta_3 \text{Dir}_i \times \text{Coord}_i)$ —using briefing slider values (pure treatment only). Direction $\in [0, 1]$ is increasing in regime strength; Coordination $\in [0, 1]$ is decreasing (high = weaker regime). Because both are monotone logistic transforms of the same signal, they are near-collinear ($\rho \approx -1$); individual coefficients reflect the partial effect *conditional on the other*, not the marginal effect. Column (3): partial effect of elicited belief ($b_i \in [0, 100]$, stated success probability) on action, controlling for z-score. “Treat: X” rows are treatment intercept dummies (base category: communication), not interactions with belief. **Caveat:** beliefs are elicited *after* the join/stay decision (post-decision), so the belief coefficient reflects association rather than a causal pathway; post-decision rationalization may inflate the correlation. The very high pseudo- R^2 (0.975) and large intercept indicate near-deterministic prediction consistent with quasi-separation—belief almost perfectly predicts action, as expected when belief is stated *post hoc*. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A15: Logistic fit parameters by model and treatment. Fitted form: $P(\text{join} | \theta) = 1/(1 + e^{b_0 + b_1 \theta})$, where $b_1 > 0$ implies a *decreasing* join curve. $\hat{\theta}^*$ is the estimated cutoff ($-b_0/b_1$); $\beta \equiv b_1$ is the steepness parameter.

Model	Pure		Communication	
	$\hat{\theta}^*$ (SE)	β (SE)	$\hat{\theta}^*$ (SE)	β (SE)
Mistral Small Creative	-0.32 (0.02)	+2.15 (0.08)	-0.23 (0.02)	+2.57 (0.11)
Llama 3.3 70B	+0.02 (0.04)	+2.83 (0.36)	-0.06 (0.04)	+3.63 (0.52)
Qwen3 30B	+0.10 (0.06)	+1.62 (0.16)	-0.03 (0.04)	+2.94 (0.36)
GPT-OSS 120B	-0.25 (0.04)	+2.06 (0.15)	-0.15 (0.03)	+3.03 (0.26)
Qwen3 235B	-0.22 (0.03)	+2.11 (0.15)	-0.22 (0.03)	+2.62 (0.23)
Trinity Large	+0.08 (0.05)	+1.34 (0.11)	-0.02 (0.04)	+2.23 (0.23)
MiniMax M2-Her	-0.17 (0.09)	+0.61 (0.06)	-0.07 (0.09)	+0.66 (0.06)

Notes: Fitted form: $P(\text{join} | \theta) = 1/(1 + e^{b_0 + b_1 \theta})$, so $b_1 > 0$ corresponds to a decreasing join curve (the expected direction). $\hat{\theta}^* = -b_0/b_1$: estimated cutoff. $\beta \equiv b_1$: logistic steepness (larger = sharper threshold). Standard errors from the covariance matrix of the nonlinear fit; cutoff SE by delta method.




Figure A8: Cross-generator robustness. Three text formats produce indistinguishable sigmoids.

Table A16: Cross-generator robustness. Three text rendering styles (baseline, diplomatic cable, journalistic wire) use identical slider functions and evidence items; only prose formatting differs. The Pearson $r(J, A(\theta))$ and logistic cutoff are virtually identical across generators.

Model	Generator	N	Mean join	$r(J, A(\theta))$	Cutoff $\hat{\theta}^*$
Mistral Small Creative	Baseline	100	0.402	+0.77	-0.070
Mistral Small Creative	Cable	100	0.406	+0.76	-0.057
Mistral Small Creative	Journalistic	100	0.403	+0.75	-0.065
Llama 3.3 70B	Baseline	100	0.354	+0.75	-0.249
Llama 3.3 70B	Cable	100	0.362	+0.75	-0.226
Llama 3.3 70B	Journalistic	100	0.360	+0.76	-0.233

Notes: Information design θ -grid. Three text renderers (baseline intelligence briefing, diplomatic cable, journalistic wire) use identical slider functions and evidence items; only prose formatting differs.

Table A17: Agent-level belief correlations (primary model: Mistral Small Creative).

Treatment	N	r_{post}	$r_{\text{b,d}}$	r_{partial}	Mean belief
Pure	10000	+0.78	+0.84	+0.64	0.444
Communication	5000	+0.79	+0.83	+0.56	0.444
Surveillance	5000	+0.78	+0.73	+0.47	0.436

Notes: r_{post} : posterior vs. stated belief. $r_{\text{b,d}}$: belief vs. decision. r_{partial} : belief-decision controlling for z-score.

Table A18: Elicited punishment risk (0–10 scale). Agents rate expected regime punishment after their JOIN/STAY decision. “JOIN” and “STAY” columns show the mean rating conditional on the agent’s own decision.

Model	Condition	N	Mean risk	Risk JOIN	Risk STAY
Mistral	Pure	2500	8.1	8.0	8.2
Mistral	Comm	2500	8.1	8.0	8.1
Mistral	Surveillance	2500	8.1	8.1	8.1
Llama 70B	Pure	2500	8.0	8.0	7.9
Llama 70B	Comm	2500	7.9	8.0	7.9
Llama 70B	Surveillance	2500	7.9	8.0	7.9

Notes: Agent-level elicitation on a 0–10 scale. Reported risk is the mean rating conditional on the agent’s own JOIN or STAY decision.




Figure A9: Cross-model signal monotonicity. Points report $r(J, A(\theta))$ under pure and communication.




Figure A10: Text baseline test. The LLM join curve is steeper than the naïve text predictor ($1 - d$), indicating processing beyond surface sentiment.

Table A19: Stakes conditions. *Top*: pipeline-centering B/C sweep—only (B, C) vary. *Bottom*: narrative stakes ladder—only a prepended header varies.

Condition / target θ^*	B	C	Header / narrative framing
<i>Payoff sweep (standard briefings, no header):</i>			
$\theta^* = 0.25$	0.33	1	—
$\theta^* = 0.33$	0.49	1	—
$\theta^* = 0.45$	0.82	1	—
$\theta^* = 0.50$	1.00	1	—
$\theta^* = 0.60$	1.50	1	—
$\theta^* = 0.67$	2.03	1	—
$\theta^* = 0.75$	3.00	1	—
<i>Narrative cost ladder ($B=C=1$; header prepended):</i>			
Very high cost	1	1	Disappearances, families targeted
Moderate high	1	1	Prison sentences, job loss
Neutral	1	1	Consequences uncertain
Moderate low	1	1	Brief questioning, released
Very low cost	1	1	Guaranteed asylum, zero cost
<i>Controls:</i>			
Placebo	1	1	Weather information (irrelevant)
High cost v2	1	1	Paraphrase of very high cost
Low cost v2	1	1	Paraphrase of very low cost

Notes: Payoff sweep: $B = \theta^*/(1-\theta^*)$, $C=1$. Z-score center and θ -grid shift with θ^* , so the briefing-text distribution is approximately constant across the seven cells. Narrative ladder: $B=C=1$ in every row; only the prepended header varies. Join rates: very high cost 21.3%, moderate high 46.6%, neutral 46.7%, moderate low 68.4%, very low 78.9%. Placebo 32.7%, v2 variants 27.4%/63.8%. Full headers in Appendix A.




Figure A11: Fitted cutoff $\hat{\theta}^*$ vs. theoretical θ^* across seven B/C ratios ($r = 0.997$). All conditions use standard briefings; the z-score grid shifts with θ^* , so this tests pipeline centering rather than direct payoff sensitivity.

$z = -2$ (weak regime)	$z = 0$ (ambiguous)	$z = +2$ (strong regime)
PRIVATE BRIEFING BOTTOM LINE: Signs of institutional erosion are accumulating faster than the regime can manage them. ATMOSPHERE: Coordination is happening in plain sight – meetings, signals, preparations that would have been unthinkable a month ago. EVIDENCE: - [Elite Cohesion] senior figures are giving interviews that subtly distance themselves from the leader's recent decisions - [Security Forces] reports of soldiers fraternizing with civilians more than usual at checkpoints (overwhelming convergence – every channel telling the same story independently) - [Money And Logistics] customs and trade are flowing normally at the borders – no unusual delays or inspections (overwhelming convergence – every channel telling the same story independently) - [Street Mood] jokes about the regime are being told openly in places where that would have been unthinkable - [Information Control] leaked documents are circulating faster than the regime can discredit them (overwhelming convergence – every channel telling the same story independently) - [Personal Observations] your cousin in the army hasn't called in two weeks – unusual for someone who calls every Sunday (near-unanimous confirmation from every reliable source) - [Diplomatic Signals] international shipping companies have reclassified the country as higher-risk for cargo insurance – the trend cooling off, though not yet cold - [Institutional Functioning] municipal services – garbage collection, water, permits – have become erratic in the capital UNCERTAINTIES: - Whether any of the apparent weakness is being deliberately staged to flush out opponents. - How much of the disarray is structural versus a temporary response to a specific crisis.	PRIVATE BRIEFING BOTTOM LINE: The situation could tip either way – there are credible signals pointing in both directions. ATMOSPHERE: There's a palpable sense that a critical mass is forming, even if nobody can count the numbers yet. EVIDENCE: - [Elite Cohesion] elites are publicly unified but privately positioning themselves for multiple outcomes (a couple of contacts mentioning it independently) - [Security Forces] the security forces are maintaining presence but seem to be conserving resources rather than demonstrating strength (a couple of contacts mentioning it independently) - [Money And Logistics] economic indicators are mixed – retail spending is down but construction permits are up - [Street Mood] people are frustrated but fatalistic – complaining more, but not in a way that suggests action - [Information Control] the regime is still shaping the narrative but has lost the initiative – they're responding to events rather than framing them - [Personal Observations] a well-connected friend said 'it could go either way' and seemed genuinely unsure (a couple of contacts mentioning it independently) - [Diplomatic Signals] foreign embassies are maintaining full operations but have updated contingency plans – standard prudence or genuine concern - [Institutional Functioning] the system is creaking but not collapsing – people are working around the bottlenecks rather than through them (a couple of contacts mentioning it independently) UNCERTAINTIES: - Whether the population's quiescence reflects genuine acceptance or merely the absence of a triggering event. - Whether the stability you're seeing is structural or just a temporary equilibrium that could shift quickly.	PRIVATE BRIEFING BOTTOM LINE: Current indicators point to maintained control, with internal challenges being managed rather than spiraling. ATMOSPHERE: People are aware of each other's frustrations but nobody is naming them directly. It's communication through omission. EVIDENCE: - [Elite Cohesion] the regime successfully managed a minor scandal by presenting a unified response within hours - [Security Forces] security checkpoints have been reduced, suggesting confidence rather than siege mentality (overwhelming convergence – every channel telling the same story independently) - [Money And Logistics] customs and trade are flowing normally at the borders – no unusual delays or inspections (overwhelming convergence – every channel telling the same story independently) - [Street Mood] people complain but do it like people who expect tomorrow to look much like today - [Information Control] social media monitoring appears sophisticated and well-resourced (overwhelming convergence – every channel telling the same story independently) - [Personal Observations] your cousin in the army hasn't called in two weeks – unusual for someone who calls every Sunday (near-unanimous confirmation from every reliable source) - [Diplomatic Signals] foreign students continue to enroll at local universities – their embassies haven't advised them to leave – the trend moving with a momentum that would be hard to reverse - [Institutional Functioning] municipal services are running smoothly – a mundane but reliable indicator of institutional health UNCERTAINTIES: - Whether the regime's strength is being accurately perceived by potential challengers – overestimation could be self-fulfilling. - If unrest flares, the response capacity seems intact – though no system is invulnerable to surprises.

Figure A12: Example briefings at three z-scores. *Left:* $z = -2$ (strong join signal) — evidence of elite fracturing, security friction, and capital flight. *Center:* $z = 0$ (ambiguous) — balanced mixed signals across all domains. *Right:* $z = +2$ (strong stay signal) — evidence of elite cohesion, military confidence, and economic stability. The displayed text is an unedited generator draw at seed 5150, agent 8, period 124 (aside from line wrapping).

Table A20: Strategic-stakes narrative comparative statics. High-cost framing emphasizes severe reprisals for failed action; low-cost framing emphasizes minimal consequences. The table is evidence on qualitative stakes framing, not literal numeric payoff reasoning.

Design	N	Mean join	$r(J, A(\theta))$	Cutoff $\hat{\theta}^*$ (SE)	Δ (pp)
Baseline	270	0.409	+0.88	0.39 (0.007)	—
High cost	270	0.190	+0.87	0.13 (0.010)	-21.9
Low cost	270	0.693	+0.88	0.72 (0.007)	+28.4

Notes: Primary model, information design θ -grid. Identical briefings across conditions; only the strategic-stakes header varies. Δ : change vs. baseline mean join (pp).

Table A21: Slider values at representative z-scores.

z	Direction	Clarity	Coordination
-2.0	0.17	0.98	0.77
-1.0	0.31	0.63	0.65
0.0	0.50	0.00	0.50
+1.0	0.69	0.63	0.35
+2.0	0.83	0.98	0.23

Table A22: Parse error and API failure rates by model and treatment.

Model	Treat.	N	API err	Unparse.	Combined
Mistral	Pure	1000	0.0%	0.0%	0.0%
	Comm	1000	0.0%	0.0%	0.0%
	Scr.	200	0.0%	0.0%	0.0%
	Flip	100	0.0%	0.0%	0.0%
Llama 70B	Pure	100	0.0%	0.0%	0.0%
	Comm	100	0.0%	0.0%	0.0%
	Scr.	100	0.0%	0.0%	0.0%
	Flip	100	0.0%	0.0%	0.0%
Qwen 30B	Pure	100	0.0%	0.0%	0.0%
	Comm	100	0.0%	0.0%	0.0%
	Scr.	100	0.0%	0.0%	0.0%
	Flip	100	0.0%	0.0%	0.0%
GPT-OSS	Pure	200	0.0%	1.5%	1.5%
	Comm	200	0.0%	0.3%	0.3%
	Scr.	500	0.0%	3.1%	3.1%
	Flip	500	0.0%	3.3%	3.3%
Qwen 235B	Pure	200	0.0%	0.0%	0.0%
	Comm	200	0.0%	0.0%	0.0%
	Scr.	100	0.0%	0.0%	0.0%
	Flip	100	0.0%	0.0%	0.0%
Trinity	Pure	100	9.0%	0.0%	9.0%
	Comm	100	10.0%	0.0%	10.0%
	Scr.	100	8.0%	0.0%	8.0%
	Flip	100	9.9%	0.0%	9.9%
MiniMax	Pure	100	0.0%	1.5%	1.5%
	Comm	100	0.0%	0.9%	0.9%
	Scr.	100	0.0%	1.8%	1.8%
	Flip	100	0.0%	1.0%	1.0%

Notes: Per-period averages. API error = provider-side failure (time-out, rate limit, content filter). Unparseable = valid response not classified as JOIN/STAY. Combined < 2% for five of seven models; Trinity Large has elevated API errors ($\approx 9\%$) due to content filtering.